

○○  
○○○○○○○○○○  
○○○○○○○○○○  
○○○○

○○○

# Bounding Techniques for Dynamic Partial Order Reduction

Σαχίνογλου Γιάννης

ΣΗΜΜΥ - ΕΜΠ

03112089

○○  
○○○○○○○○○○  
○○○○○○○○○○  
○○○○

○○○

## Περίληψη

○○  
○○○○○○○○○○  
○○○○○○○○○○  
○○○○

○○○

## Θέμα εργασίας



## Concurrent Computing and Problems

**Concurrent Computing:** Concurrent computing is a form of computing in which several computations are executed during overlapping time periods-concurrently-instead of sequentially (one completing before the next starts).

Potential problems include:

- Race Conditions
- Deadlocks
- Livelocks
- Resource Starvation



## Concurrency Errors

### A simple example:

```
void *divider(void* arg){  
    int x = 0;  
    return 42/x;  
}
```

Listing 1: Example of non-concurrency error

```
volatile int x = 1;  
void *divider(){  
    return 42/x;  
}  
  
void *zero(){  
    x = 0;  
}
```

Listing 2: Example of concurrency error



## Testing, Model Checking, and Verification

- Testing: For some given inputs check whether the output is correct.
- Verification: Prove formally that the output is correct.
- Model Checking: Explore all the possible states the system can be.

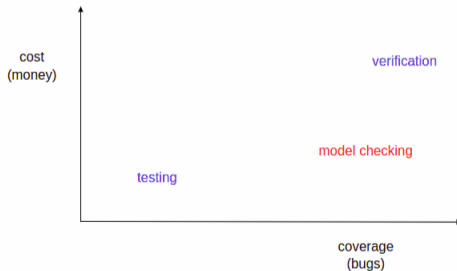


Figure: Comparing Testing, Model Checking and Verification



## The Idea of Interleaving

- We need to model our state space!
- An Interleaving represents a scheduling of the concurrent program.
- In order to find an error of a concurrent program, one must examine every possible interleaving BUT leads to state explosion.



## Stateless Model Checking and Partial Order Reduction

Partial order reduction aims to reduce the number of interleavings explored by eliminating the exploration of equivalent interleavings.

For example:




Figure: Examples of Interleavings





## Stateless Model Checking and Partial Order Reduction

- Static Partial Order Reduction: Dependencies are tracked before execution.
- Dynamic Partial Order Reduction: Dependencies are observed during runtime.



## Bounding Techniques for DPOR

- For larger programs DPOR often runs longer than developers are willing to wait.
- Bounded techniques, alleviate state-space explosion by pruning the executions that exceed a bound.
- Preemption Bounded and Delay Bounded exploration.
- Many of the concurrency bugs can be tracked even when the bound limit is set to be small.

○○  
○○○○○●○○○  
○○○○○○○○○○  
○○○

○○○

We need to introduce some basic ideas and notation!



## Vector Clocks

1. Each process experiencing an internal event, it increments its own logical clock in the vector by one.
2. Each time a process receives a message or performs an action on a shared variable, it increments its own logical clock in the vector by one and updates each element in its vector by taking the maximum of the value in its own vector clock and the value in the vector in the received message or the maximum value of all processes that share the same shared variable.

○○  
○○○○○○○○●○○  
○○○○○○○○○○○  
○○○○

○○○

## Useful Notation




## Event Dependencies

### Definition 1 (happens-before assignment)

A happens-before assignment, which assigns a unique happens-before relation  $\rightarrow_E$  to any execution sequence  $E$ , is valid if it satisfies the following properties for all execution sequences  $E$ .

1.  $\rightarrow_E$  is a partial order on  $dom(E)$ , which is included in  $<_E$ . In other words every scheduling is part of the set of all possible partial order of the program.
2. The execution steps of each process are totally ordered, i.e.  
 $\langle p, i \rangle \rightarrow_E \langle p, i + 1 \rangle$  whenever  $\langle p, i + 1 \rangle \in dom(E)$ .
3. If  $E'$  is a prefix of  $E$  then  $\rightarrow_E$  and  $\rightarrow_{E'}$  are the same on  $dom(E')$ .



## Event Dependencies

4. Any linearization  $E'$  of  $\rightarrow_E$  on  $dom(E)$  is an execution sequence which has exactly the same “happens-before” relation  $\rightarrow_{E'}$  as  $\rightarrow_E$ . This means that the relation  $\rightarrow_E$  induces a set of equivalent execution sequences, all with the same “happens-before” relation. We use  $E \simeq E'$  to denote that  $E$  and  $E'$  are linearizations of the same “happens-before” relation, and  $[E] \simeq$  to denote the equivalence class of  $E$ .
5. If  $E \simeq E'$  then  $s_{[E]} = s_{[E']}$  (i.e. two equivalent traces will lead to the same state).
6. For any sequences  $E, E'$  and  $w$ , such that  $E.w$  is an execution sequence, we have  $E \simeq E'$  if and only if  $E.w \simeq E'.w$ .

```

○○
○○○○○○○○○○
●○○○○○○○○○○
○○○○

```

```

○○○

```

## Definition 2 (Sufficient Sets)

A set of transitions is sufficient in a state  $s$  if any relevant state reachable via an enabled transition from  $s$  is also reachable from  $s$  via at least one of the transitions in the sufficient set. A search can thus explore only the transitions in the sufficient set from  $s$  because all relevant states still remain reachable. The set containing all enabled threads is trivially sufficient in  $s$ , but smaller sufficient sets enable more state space reduction.





## General form of DPOR

Explore( $\emptyset$ );

**Function** *Explore*( $E$ )

**let**  $T = Sufficient\_set(final(E))$ ;

**for all**  $t \in T$  **do**

        Explore( $E.t$ ) ;

**end**

**Algorithm 1:** General form of DPOR



## Sufficient Sets: Persistent Sets

### Definition 3 (Persistent Sets)

Let  $s$  be a state, and let  $W \subseteq E(s)$  be a set of execution sequences from  $s$ . A set  $T$  of transitions is a persistent set for  $W$  after  $s$  if for each prefix  $w$  of some sequence in  $W$ , which contains no occurrence of a transition in  $T$ , we have  $E \vdash t \Diamond w$  for each  $t \in T$ .



## Sufficient Sets: Persistent Sets

A simple example:




Figure: Construction of persistent sets



## Sufficient Sets: Source Sets

### Definition 4 ( $dom(E)$ )

The set of events-transitions happening during the scheduling of  $E$ .

### Definition 5 (Initials after an execution sequence $E.w$ , $I_{[E]}(w)$ )

For an execution sequence  $E.w$ , let  $I_{[E]}(w)$  denote the set of processes that perform events  $e$  in  $dom_{[E]}(w)$  that have no “happens-before” predecessors in  $dom_{[E]}(w)$ . More formally,  $p \in I_{[E]}(w)$  if  $p \in w$  and there is no other event  $e \in dom_{[E]}(w)$  with  $e \rightarrow_{E.w} next_{[E]}(p)$ .

By relaxing the definition of Initials we can get the definition of Weak Initials,  $WI$ .

### Definition 6 (Weak Initials after an execution sequence $E.w$ , $WI_{[E]}(w)$ )

For an execution sequence  $E.w$ , let  $WI_{[E]}(w)$  denote the union of  $I_{[E]}(w)$  and the set of processes that perform events  $p$  such that  $p \in enabled(s_{[E]})$ .



## Sufficient Sets: Source Sets



## Sufficient Sets: Source Sets

### Definition 7 (Source Sets)

Let  $E$  be an execution sequence, and let  $W$  be a set of sequences, such that  $E.w$  is an execution sequence for each  $w \in W$ . A set  $T$  of processes is a source set for  $W$  after  $E$  if for each  $w \in W$  we have  $WI_{[E]}(w) \cap T = \emptyset$ .



## Source Sets

An example:

We don't need to add  $r$  since  $q$  already belongs to source set.




Figure: Construction of Source Sets



## Further Optimizations: Sleep Sets

The idea behind Sleep Set Optimization:

- Assume that the search explores transition  $t$  from state  $s$ , backtracks  $t$ , then explores  $t_0$  from  $s$  instead. Unless the search explores a transition that is dependent with  $t$ , no states are reachable via  $t_0$  that were not already reachable via  $t$  from  $s$ . Thus,  $t$  “sleeps” unless a dependent transition is explored.





## Sleep Sets

Sleeps sets in action (Using Persistent Sets):




Figure: Example of Sleep Set Optimization



## Bounded Dynamic Partial Order Reduction General Form

Given a bound evaluation function  $B_v$  and a bound  $c$ :

**Result:** Explore the whole statespace

Explore( $\emptyset$ );

**Function** *Explore*( $E$ )

$T = \text{Sufficient\_set}(\text{final}(E))$

**for** all  $t \in T$  **do**

**if**  $B_v(E.t) \leq c$  **then**

            Explore( $E.t$ )

**end**

**end**

**Algorithm 2:** Bounded-DPOR



## Preemption Bounded Search

### Definition 8 (Preemption bound)

$$P_b(\emptyset) = 0$$

$$P_b(E.t) =$$

$$\begin{cases} P_b(E) + 1 & \text{if } t.tid = last(E).tid \text{ and } last(E).tid \in enabled(final(E)) \\ P_b(E) & \text{otherwise} \end{cases}$$

```

○○
○○○○○○○○○○
○○○○○○○○○○
○○●○

```

```

○○○

```

## Definition 9 ( $ext(s, t)$ )

Given a state  $s = final(E)$  and a transition  $t \in enabled(s)$ ,  $ext(s, t)$  returns the unique sequence of transitions  $\beta$  from  $s$  such that

1.  $\forall i \in dom(\beta) : \beta_i.tid = t.tid$
2.  $t.tid \notin enabled(final(E.\beta))$



## Preemption Bounded Persistent Sets

### Definition 10 (Preemption Bounded Persistent Set)

A set  $T \subseteq \mathcal{T}$  of transitions enabled in a state  $s = final(E)$  is preemption-bound persistent in  $s$  iff for all nonempty sequences  $a$  of transitions from  $s$  in  $A_G(P_b, c)$  such that  $\forall i \in dom(a), a_i \notin T$  for all  $t \in T$ ,

1.  $Pb(E.t) \leq Pb(E.a_1)$
2. if  $Pb(E.t) < Pb(E.a_1)$ , then  $t \leftrightarrow last(a)$  and  $t \leftrightarrow next(final(E.a), last(a).tid)$
3. if  $Pb(E.t) = Pb(E.a_1)$ , then  $ext(s, t) \leftrightarrow last(a)$  and  $ext(s, t) \leftrightarrow next(final(E.a), last(a).tid)$

```

○○
○○○○○○○○○○
○○○○○○○○○○
○○○○

```

```

○○○

```

## Source-DPOR

```

Explore( $\langle \rangle, \emptyset$ );
Function Explore( $E, Sleep$ )
  if  $\exists p \in (enabled(s_{[E]}) \setminus Sleep)$  then
     $backtrack(E) := p$ ;
    while  $\exists p \in (backtrack(E) \setminus Sleep)$  do
      foreach  $e \in dom(E)$  such that  $e \lesssim_{E,p} next_{[E]}(p)$  do
        let  $E' = pre(E, e)$ ;
        let  $u = notdep(e, E).p$ ;
        if  $I_{E'}(u) \cap backtrack(E') = \emptyset$  then
          | add some  $q' \in I_{[E']}(u)$  to  $backtrack(E')$ ;
        end
      end
      let  $Sleep' := \{q \in Sleep \mid E \models p \diamond q\}$ ;
       $Explore(E.p, Sleep')$ ;
      add  $p$  to  $Sleep$ ;
    end
  end

```

**Algorithm 3:** Source-DPOR Algorithm

```

○○
○○○○○○○○○○
○○○○○○○○○○
○○○○

```

```

○○○

```

## DPOR using Clock Vectors (Classic-DPOR)

**Function** *Explore*( $E, C$ )

```

let  $s := \text{last}(E)$ ;
for all process  $p$  do
  if  $\exists i = \max(\{i \in \text{dom}(E) \mid E_i \text{ is dependent and may be co-enabled with } \text{next}(s, p) \text{ and } i \notin C(p)(\text{proc}(E_i))\})$  then
    if  $p \in \text{enabled}(\text{pre}(E, i))$  then
       $\text{add } p \text{ to } \text{backtrack}(\text{pre}(E, i))$  ;
    else
       $\text{add } \text{enabled}(\text{pre}(E, i)) \text{ to } \text{backtrack}(\text{pre}(E, i))$  ;
    end
  end
end
if  $\exists p \in \text{enabled}(s)$  then
   $\text{backtrack}(s) := p$  ;
  let  $\text{done} = \emptyset$ ;
  while  $\exists p \in (\text{backtrack}(s) \setminus \text{done})$  do
     $\text{add } p \text{ to } \text{done}$  ;
    let  $t = \text{next}(s, p)$ ;
    let  $E' = E.t$ ;
    let  $cu = \max\{C(i) \mid i \in 1..|S| \text{ and } E_i \text{ dependent with } t\}$ ;
    let  $cu2 = cu[p := |E'|]$ ;
    let  $C' = C[p := cu2, |E'| := cu2]$ ;
    Explore( $E', C'$ ) ;
  end
end

```

**Algorithm 4:** DPOR using Clock Vectors (Classic-DPOR)



## Source-DPOR vs Classic-DPOR

### Similarities:

1. Consist of the same phases i.e., race detection and exploration
2. Both rely on Vector Clocks.

### Differences:

1. Classic-DPOR “eager” i.e., adds more dependencies before scheduling.
2. Source-DPOR “lazy” i.e., adds branches after scheduling and thus avoids redundant additions.





## Nidhugg-DPOR

```

Explore( $\langle \rangle, \emptyset$ );
Function Explore( $E, Sleep$ )
  if  $\exists p \in (enabled(s_{[E]}) \setminus Sleep)$  then
    backtrack( $E$ ) :=  $p$  ;
    while  $\exists p \in (backtrack(E) \setminus Sleep)$  do
      foreach  $e \in dom(E)$  such that  $e \lesssim_{E.p} next_{[E]}(p)$  do
        let  $E' = pre(E, e)$ ;
        let  $u = notdep(e, E).p$ ;
        let  $CI = \{i \in I_{E'}(u) \mid i \rightarrow p\}$ ;
        if  $CI \cap backtrack(E') = \emptyset$  then
          if  $CI \neq \emptyset$  then
            | add some  $q' \in CI$  to backtrack( $E'$ ) ;
          end
          else
            | add some  $q' I_{E'}(u)$  to backtrack( $E'$ )
          end
        end
      end
    let  $Sleep' := \{q \in Sleep \mid E \models p \triangleleft q\}$  ;
    Explore( $E.p, Sleep'$ ) ;
    add  $p$  to Sleep ;
  end
end

```


Algorithm 5: Nidhugg-DPOR



## Correctness of Nidhugg-DPOR

Case 1: At least one process contains a write command. We know that the two processes will be inverted at some point. Since Nidhugg-DPOR ignores weak initials it will branch both processes. In Source-DPOR only one of the two processes should be branched since they share the same initials. However, in Nidhugg-DPOR this is not true since the *CI* set does not contain steps from the other process.

When *r* is added *q* is not considered since it does not belong to *CI*  
in contrast to *I* set of Source DPOR




**Figure:** Construction of persistent sets in Nidhugg when there is a write process



## Correctness of Nidhugg-DPOR

Case 2: Both processes are read operations. Since we do not calculate  $I$  but  $CI$  the first read operation will not be considered as it does not happen before the second read operation and as result both processes will be added to *backtrack*. We notice that by calculating the  $CI$  set when the race between  $p$  and  $r$  is detected  $q$  process will be ignored and, thus,  $r$  will be added as a branch.

When  $r$  is added  $q$  is not considered since it does not belong to  $CI$  in contrast to  $I$  set of Source DPOR



**Figure:** Construction of persistent sets in Nidhugg when both are read processes



## Naive-BPOR

```

Explore( $\langle \rangle, \emptyset, b$ );
Function Explore( $E, Sleep, b$ )
  if  $\exists p \in (enabled(s_{[E]}) \setminus Sleep)$  such that  $B_v(E.p) \leq b$  then
    backtrack( $E$ ) :=  $p$  ;
    while  $\exists p \in (backtrack(E) \setminus Sleep)$  and  $B_v(E.p) \leq b$  do
      foreach  $e \in dom(E)$  such that  $e \lesssim_{E.p} next_{[E]}(p)$  do
        let  $E' = pre(E, e)$ ;
        let  $u = notdep(e, E).p$ ;
        if  $I_{E'}(u) \cap backtrack(E') = \emptyset$  then
          add some  $q' \in I_{[E']}(u)$  to backtrack( $E'$ ) ;
        end
      end
      let  $Sleep' := \{q \in Sleep \mid E \models p \Diamond q\}$ ;
      Explore( $E.p, Sleep', b$ ) ;
      add  $p$  to  $Sleep$  ;
    end
  end

```

**Algorithm 6:** Naive-BPOR



## Example execution of Naive-BPOR

A Naive-BPOR execution example and the problem with it.




Figure: Naive-BPOR for bound=0

○○

○○○○○○○○○○

○○○○○○○○○○

○○○○

○○○

# Classic-BPOR

```

Function Explore(E)
  let s := last(E);
  for all process p do
    for all process q ≠ p do
      if ∃ i = max({i ∈ dom(E) | Ei is dependent and may be co-enabled
        with next(s, p) and Ei.tid = q then
        if p ∈ enabled(pre(E, i)) then
          add p to backtrack(pre(E, i)) ;
        else
          add enabled(pre(E, i)) to backtrack(pre(E, i)) ;
        end
      if j = max({j ∈ dom(E) | j = 0 or Sj-1.tid ≠ Sj.tid and j < i})
        then
          if p ∈ enabled(pre(E, i)) then
            add p to backtrack(pre(E, i)) ;
          else
            add enabled(pre(E, i)) to backtrack(pre(E, i)) ;
          end
        end
      end
    end
  end
  if p ∈ enabled(s) then
    add p to backtrack(s) ;
  end
  else
    add any u ∈ enabled(s) to backtrack(s) ;
  end
  let visited = ∅;
  while ∃ u ∈ (enabled(s) ∩ backtrack(s) \ visited) do
    add u to visited ;
    if (Bc(S.next(s, u)) ≤ c) then
      Explore(S.next(s, u)) ;
    end
  end

```

Algorithm 7: BPOR

```

○○
○○○○○○○○○○
○○○○○○○○○○
○○○○

```

```

○○○

```

## Nidhugg-BPOR

```

Explore( $\emptyset, \emptyset, b$ );
Function Explore( $E, Sleep, b$ )
  backtrack( $E$ ) :=  $p$ ;
  while  $\exists p \in (enabled(s_{[E]}) \setminus Sleep)$  and  $B_v(E, p) \leq b$  then
    foreach  $e \in dom(E)$  such that  $e \lesssim_{E, p} next_{[E]}(p)$  do
      let  $E' = pre(E, e)$ ;
      let  $u = notdep(e, E).p$ ;
      let  $CI = \{i \in I_{E'}(u) \mid i \rightarrow p\}$ ;
      if  $CI \cap backtrack(E') = \emptyset$  then
        if  $CI \neq \emptyset$  then
          add some  $q' \in CI$  to  $backtrack(E')$ ;
        end
        else
          add some  $q' \in I_{E'}(u)$  to  $backtrack(E')$ ;
        end
      end
      let  $E'' = pre\_block(e, E)$ ;
      let  $u = notdep(e, E).p$ ;
      let  $CI = \{i \in I_{E''}(u) \mid i \rightarrow p\}$ ;
      if  $CI \cap backtrack(E') = \emptyset$  then
        if  $CI \neq \emptyset$  then
          add some  $q' \in CI$  to  $backtrack(E')$ ;
        end
        else
          add some  $c(q') \in I_{E''}(u)$  to  $backtrack(E'')$ ;
        end
      end
    end
  end
  let  $Sleep' := \{q \in Sleep \mid E \models p \hat{\Delta} q\}$ ;
  Explore( $E.p, Sleep'$ );
  if  $p$  is not conservative then
    add  $p$  to  $Sleep$ ;
  end
end

```

Algorithm 8: Nidhugg-BPOR

○○  
○○○○○○○○○○  
○○○○○○○○○○  
○○○○

○○○

## The main question

Can we use source sets instead of persistent sets in order implement BPOR?





```
○○
○○○○○○○○○○
○○○○○○○○○○
○○○○
```

```
○○○
```

## A Correct Approach

We should use Source Sets for non-conservative branches and persistent sets for conservative branches.

```

OO
OOOOOOOOOO
OOOOOOOOOO
OOOO

```

```

OOO

```

## Source-BPOR

```

Explore( $\emptyset, \emptyset, b$ );
Function Explore( $E, Sleep, b$ )
  if  $\exists p \in ((enabled(s_{[E]}) \backslash Sleep) \text{ and } B_c(E.p) \leq b)$  then
    backtrack( $E$ ) :=  $p$ ;
    while  $\exists p \in (backtrack(E) \backslash Sleep \text{ and } B_c(E.p) \leq b)$  do
      foreach  $e \in dom(E)$  such that  $e \lesssim_{E.p} next_{[E]}(p)$  do
        let  $E' = pre(E, e)$ ;
        let  $u = notdep(e, E).p$ ;
        if  $I_{E'}(u) \cap backtrack(E') = \emptyset$  then
          | add some  $q' \in I_{E'}(u)$  to backtrack( $E'$ );
        end
        let  $E'' = pre\_block(e, E)$ ;
        let  $u = notdep(e, E).p$ ;
        let  $CI = \{i \in I_{E''}(u) \mid i \rightarrow p\}$ ;
        if  $CI \cap backtrack(E') = \emptyset$  then
          if  $CI \neq \emptyset$  then
            | add some  $q' \in CI$  to backtrack( $E'$ );
          end
          else
            | add some  $c(q') \in I_{E''}(u)$  to backtrack( $E''$ );
          end
        end
      end
    end
    let  $Sleep' := \{q \in Sleep \mid E \models p \Diamond q\}$ ;
    Explore( $E.p, Sleep'$ );
    if  $p$  is not conservative then
      | add  $p$  to Sleep;
    end
  end
end

```

Algorithm 9: Source-BPOR

```
○○
○○○○○○○○○○
○○○○○○○○○○
○○○
```

```
○○○
```

## Nidhugg-BPOR vs Source-BPOR

### Similarities:

- Same structure.

### Differences:

- Source-BPOR relies on Source Sets for the addition of non-conservative branches while Nidhugg-BPOR relies on persistent sets.

The usage of conservative branches leads to explosion of the state space:

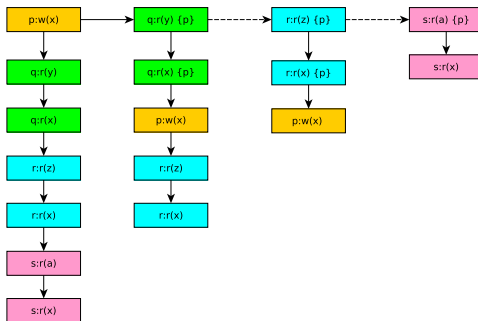


Figure: writer-3-readers explosion



Sleep Sets are no longer that useful:




Figure: Sleep set contradiction



## Concluding Remarks

The Performance - Soundness Tradeoff

○○  
○○○○○○○○○○  
○○○○○○○○○○  
○○○○

○○○

## The Nidhugg Flow Chart




Figure: Nidhugg's Flow Chart



```
○○
○○○○○○○○○○
○○○○○○○○○○
○○○○
```

```
○○○
```

The implementation mainly is focused, as expected, on `see_events()` and `add_branches()`



## Nidhugg-DPOR Evaluation

### Evaluation of Nidhugg-DPOR on Synthetic Tests




Figure: writer-N-readers

Test case	Traces for Source-DPOR	Traces for Classic-DPOR
account.c	6	7
lazy.c	6	7
micro.c	52495	53084
lastzero.c	97	97
lastzeromod.ll	13	17
indexer0.c	8	8
indexermod.c	120	226

Table: Source-DPOR vs Nidhugg-DPOR for Synthetic tests



## Evaluation of Nidhugg-BPOR on Synthetic Tests




Figure: writer-N-readers bounded

Technique:	Naïve-BPOR			Nidhugg-BPOR			Source-BPOR		
Bound:	0	1	2	0	1	2	0	1	2
account.c	1	1	4	6	27	42	6	27	42
lazy.c	1	1	4	6	27	42	6	27	42
micro.c	1	1	10	6	93	886	6	93	886
lastzero.c	1	2	5	252	2444	10614	252	2444	10614
lastzeromod.ll	1	1	6	64	290	651	64	290	651
indexer0.c	1	4	1	2	8	14	2	8	14
indexermod.c	1	1	5	120	1320	7920	120	1320	7920

Table: Traces for various bound limits

```

OO
OOOOOOOOOO
OOOOOOOOOO
OOOO

```

```

OOO

```

## Evaluation of BPOR on RCU

Read-Copy-Update (RCU): Read-copy update (RCU) is a synchronization mechanism that was added to the Linux kernel in October of 2002.

Let's start with a small bound...

ver:	3.0						3.19						4.9.6					
method:	Naive-BPOR			Classic-BPOR			Naive-BPOR			Classic-BPOR			Naive-BPOR			Classic-BPOR		
	traces	time	error	traces	time	error	traces	time	error	traces	time	error	traces	time	error	traces	time	error
-	3	0.2	NF	44	0.72	NF	2	0.32	NF	28	0.76	NF	2	0.61	NF	24	1.21	NF
-DFORCE_FAILURE_1	3	0.2	NF	44	0.72	NF	2	0.32	NF	28	0.76	NF	2	0.6	NF	24	1.21	NF
-DFORCE_FAILURE_3	3	0.2	NF	44	0.72	NF	2	0.32	NF	33	1.06	NF	2	0.61	NF	41	2.11	NF
-DFORCE_FAILURE_5	3	0.2	NF	44	0.71	NF	2	0.31	NF	18	0.55	NF	2	0.6	NF	16	0.93	NF
-DLIVENESS_CHECK_1	3	0.2	NF	44	0.72	NF	2	0.32	NF	28	0.74	NF	2	0.61	NF	24	1.19	NF
-DLIVENESS_CHECK_2	3	0.2	NF	52	0.84	NF	2	0.32	NF	28	0.73	NF	2	0.6	NF	24	1.2	NF
-DLIVENESS_CHECK_3	3	0.2	NF	44	0.71	NF	2	0.31	NF	28	0.75	NF	2	0.6	NF	24	1.19	NF

Table: RCU results for bound  $b = 1$

```

○○
○○○○○○○○○○
○○○○○○○○○○
○○○○

```

```

○○○

```

## Evaluation of BPOR on RCU

Let's increase the bound...

ver:	3.0						3.19						4.9.6					
method:	Naive-BPOR			Classic-BPOR			Naive-BPOR			Classic-BPOR			Naive-BPOR			Classic-BPOR		
	traces	time	error	traces	time	error	traces	time	error	traces	time	error	traces	time	error	traces	time	error
-	50	1.18	NF	5634	88.78	NF	10	0.49	NF	2083	60.48	NF	10	0.89	NF	2469	122.71	NF
-DFORCE FAILURE 1	50	1.06	NF	275	4.2	F	10	0.49	NF	182	5.51	F	10	0.89	NF	300	15.42	F
-DFORCE FAILURE 3	50	1.05	NF	1627	23.09	NF	15	0.72	NF	100000	0.0	NF	15	1.2	NF	100000	0.0	NF
-DFORCE FAILURE 5	49	1.05	NF	4155	59.47	NF	9	0.45	NF	60	2.34	F	9	0.81	NF	60	3.92	F
-DLIVENESS CHECK 1	48	1.04	NF	1493	21.19	NF	10	0.5	NF	517	10.66	NF	10	0.88	NF	404	13.58	NF
-DLIVENESS CHECK 2	61	1.28	NF	2105	30.5	NF	10	0.5	NF	517	10.61	NF	10	0.88	NF	582	20.28	NF
-DLIVENESS CHECK 3	49	1.04	NF	1788	24.98	NF	10	0.5	NF	655	14.04	NF	10	0.88	NF	506	17.32	NF

Table: RCU results for bound  $b = 4$

```

○○
○○○○○○○○○○
○○○○○○○○○○
○○○

```

```

○○○

```

## Evaluation of BPOR on RCU

What did we achieve?

ver:	3.0						3.19						4.9.6					
method:	Source-DPOR			Classic-BPOR			Source-DPOR			Classic-BPOR			Source-DPOR			Classic-BPOR		
	traces	time	bound	traces	time	bound	traces	time	bound	traces	time	bound	traces	time	bound	traces	time	bound
-DFORCE FAILURE 1	247	3.81	F	275	4.2	4	515	16.88	F	182	5.51	4	861	45.69	F	300	15.42	4
-DFORCE FAILURE 3	2372	33.42	NF				17094	626.4	F	201	7.03	2	15349	883.98	F	258	14.24	2
-DFORCE FAILURE 5	12426	178.8	NF				118	3.99	F	60	2.34	4	112	6.34	F	60	3.92	4

**Table:** Comparison between DPOR and BPOR with the bug



## Equivalence of Source-BPOR with Nidhugg-BPOR

### Equivalence Case1:




Figure: Source-BPOR and Nidhugg-BPOR equivalence Case 1

## Equivalence of Source-BPOR with Nidhugg-BPOR

Equivalence Case2:




Figure: Source-BPOR and Nidhugg-BPOR equivalence Case 2



```

○○
○○○○○○○○○○
○○○○○○○○○○
○○○○

```

```

○○○

```

## Motivation

Some preemption-switches can be easily avoided. For example:




Figure: An example of avoidable preemption-switch

```

○○
○○○○○○○○○○
○○○○○○○○○○
○○○

```

```

○○○

```

## Alternating the General Form of BPOR

What if calculate something more than the preemption-bound?

**Result:** Explore the whole state space within the bound

Explore( $\emptyset$ );

**Function** *Explore*(*S*)

```

| T = Sufficient_set(final(S)) for all t ∈ T do
|   if  $\min\{B_v(\overline{[S.t]})\} \leq c$  then
|     Explore(S.t)
|   end
| end

```

**Algorithm 10:** General form of the BPOR without branch addition

```

○○
○○○○○○○○○○
○○○○○○○○○○
○○○○

```

```

○○○

```

## Construct a Block-Graph using Partial Order Reduction

**Function** *AddBlock(block, graph)*

```

if previous block of the same thread was not blocked then
    | increase the weigh of the edges coming from the previous block to 1 ;
end
for each thread t do
    list:= preceding blocks t;
    for l in reversed(list) do
        if  $l \leftrightarrow block$  then
            | add edge from block to l with weight 0 ;
            if l is not last then
                | add edge from l to block with weight 1 ;
            end
        else
            | add edge from l to block with weight 0 ;
        end
    end
    if  $l \rightarrow block$  then
        if l is not last then
            | add edge from l to block with weight 1 ;
        end
        else
            | add edge from l to block with weight 0 ;
        end
        break ;
    end
end
end

```

**Algorithm 11:** Adding a new block to the dependencies' graph

```

○○
○○○○○○○○○○
○○○○○○○○○○
○○○○

```

```

○○○

```

## Applying the Graph construction




Figure: Graph example

```

○○
○○○○○○○○○○
○○○○○○○○○○
○○○○○○○○○○
○○○○

```

```

○○○

```

## Introducing Lazy-BPOR

```

let  $G =: \emptyset$ ;
Explore( $\langle \rangle, \emptyset, G, b$ );
Function Explore( $E, Sleep, G, b$ )
  if  $\exists p \in (enabled(s_{[E]}) \setminus Sleep)$  such that  $B_v(E.p) \leq b$  then
    backtrack( $E$ ) :=  $p$  ;
    while  $\exists p \in (backtrack(E) \setminus Sleep)$  do
      foreach  $e \in dom(E)$  such that  $e \lesssim_{E.p} next_{[E]}(p)$  do
        let  $E' = pre(E, e)$ ;
        let  $u = notdep(e, E).p$ ;
        if  $I_{E'}(u) \cap backtrack(E') = \emptyset$  then
          | add some  $q' \in I_{[E']}(u)$  to  $backtrack(E')$  ;
        end
      end
      let  $Sleep' := \{q \in Sleep \mid E \models p \Diamond q\}$ ;
      if  $p$  creates a new block then
        | let  $block = last\_block(E)$ ;
        | let  $G' = add\_block(block, G)$ ;
      end
      if
         $min\{Ham\_path(G') \text{ which compensate with all happens-before relations of } E\} \leq b$ 
      then
        | Explore( $E.p, Sleep, G', b$ ) ;
        | add  $p$  to  $Sleep$  ;
      end
    end
  end
end

```

Algorithm 12: Lazy-BPOR



## Evaluation of Lazy-BPOR

Evaluation on Synthetic Tests:




Figure: writer-N-readers bounded by the first estimation algorithm

Technique:	Naive-BPOR			Lazy-BPOR			Nidhugg-BPOR		
	0	1	2	0	1	2	0	1	2
account.c	1	1	4	6	6	6	6	27	42
lazy.c	1	1	4	6	6	6	6	27	42
micro.c	1	1	10	60	805	4362	6	93	886
lastzero.c	1	2	5	97	97	97	252	2444	10610
lastzeromod.ll	1	1	6	13	13	13	64	290	651
indexer0.c	1	4	1	4	8	8	2	8	14
indexermod.c	1	1	5	120	120	120	120	1320	7920

Table: Traces for the first estimation algorithm for various bound limits

[illegible]

**Table:** Comparison between DPOR and Lazy-BPOR without the bug

### Evaluation on RCU (Nidhugg-BPOR vs Lazy-BPOR):

ver	3.0						3.19						4.3						4.7						4.9.6					
method	Ndhugge-BPOR			Lazy-BPOR			Ndhugge-BPOR			Lazy-BPOR			Ndhugge-BPOR			Lazy-BPOR			Ndhugge-BPOR			Lazy-BPOR			Ndhugge-BPOR			Lazy-BPOR		
	time	traces	bound	time	traces	bound	time	traces	bound	time	traces	bound	time	traces	bound	time	traces	bound	time	traces	bound	time	traces	bound	time	traces	bound	time	traces	bound
- ASSERT 0	2.85	183	3	2.79	184	2	2.86	186	3	4.06	73	2	12.59	126	3	8.57	85	2	5.26	118	3	6.26	75	2	5.91	128	3	9.44	85	2
- DPOFCE FAILURE 0	2.85	275	4	2.84	241	3	3.62	352	4	6.48	123	3	17.73	163	4	6.26	220	4	11.72	123	3	12.03	300	4	28.04	153	3	28.04	153	3
- DPOFCE FAILURE 1	0.35	8	1	0.35	4	1	0.54	4	1	0.52	3	0	0.73	5	1	0.71	7	0	0.90	5	1	0.87	0	0	0.95	5	1	0.9	3	0
- DPOFCE FAILURE 2	0	0	0	0	0	0	6.40	201	2	54.62	260	1	12.11	258	2	103.89	233	1	12.96	258	2	107.1	235	1	12.84	258	2	111.37	233	1
- DPOFCE FAILURE 3	0.91	47	2	1.38	51	2	1.78	41	2	2.1	34	1	1.89	21	2	1.79	14	1	2.3	24	2	2.27	17	1	2.39	24	2	2.34	17	1
- DPOFCE FAILURE 4	0.95	47	2	1.38	51	2	1.78	40	4	2.1	32	1	1.82	6	4	1.58	52	3	1.47	60	4	1.64	52	3	1.61	60	4	1.66	52	3
- DPOFCE FAILURE 5	0.95	47	2	1.38	51	2	1.78	41	2	2.1	34	1	1.89	21	2	1.79	14	1	2.3	24	2	2.27	17	1	2.39	24	2	2.34	17	1

**Table:** Comparison between BPOR and Lazy-BPOR



```

○○
○○○○○○○○○○
○○○○○○○○○○
○○○○

```

```

○○○

```

## Conclusion:

- It is possible to explore a preemption-bounded state space without the addition of conservative branches.
- It provides an upper bound for the number of traces explored in BPOR no matter the bound. In fact the number of traces explored by Lazy-BPOR at worst case equal to the number of traces explored by the unbounded DPOR. This is true since no conservative branches are added.
- The most important is that provides a reduction of the preemption-bounded search to a well known graph problem where many heuristics can be applied in order to expedite the calculation of the minimum hamiltonian path.

○○  
○○○○○○○○○○  
○○○○○○○○○○  
○○○○

○○○